

Supplementary Material for “Evaluating infectious disease forecasts with allocation scoring rules”

Aaron Gerding, Nicholas G. Reich, Benjamin Rogers, Evan L. Ray

December 21, 2023

Contents

1	Introduction	1
2	Proper scoring rules	1
2.1	The allocation score is proper	2
3	Expected shortages	2
4	Allocation Bayes acts as vectors of marginal quantiles.	3
5	Numerical computation of allocation Bayes acts	4
6	Computing allocations from finite quantile forecast representations	4
7	Propriety of parametric approximation	5
	□ 2 actual references to supplement (re exponential quantiles and solving the allocation problem)	
	□ the statement that “we computed the allocation score for the 14 day-ahead forecast” which needs some answer to “how?”	
	□ use of “proper” a number of times	

1 Introduction

ELR:[This section needs to be updated.]

We briefly address some technical and methodological points in the main text. We begin in section ?? by formalizing the concept of a *shortage* of resources and giving some key results about expected resource shortages under a distribution characterizing uncertainty about (future) levels of resource need. Resource shortages play a central role in the decision-making problems that give rise to the quantile loss and the allocation score, which we discuss in sections ?? and 4 respectively. Section 5 gives details on the numerical methods that we use to calculate allocation scores, including some special considerations for settings where forecasts are represented by a finite collection of predictive quantiles, such as the application to forecasts of hospitalizations due to COVID-19 in section 3 of the article.

2 Proper scoring rules

In decision theory, a loss function l is used to formalize a decision problem by assigning numerical value $l(x, y)$ to the *result* of taking an *action* x in preparation for an *outcome* y . A *scoring rule* S is a loss function for a decision problem where the action is a probabilistic forecast F of the outcome y (or the statement of F by a forecaster). We refer to the realized loss $S(F, y)$ as the *score* of F at y .

Probabilistic forecasts can be seen as a unique kind of action in that they can be used to generate their own (simulated) outcome data, against which they can be scored using S . A probabilistic forecast F is thus committed to the “self-assessment” $\mathbb{E}_F[S(F, Y)] := \mathbb{E}[S(F, Y^F)]$, where $Y^F \sim F$ is the random variable defined by sampling from F , as well to an assessment $\mathbb{E}_F[S(G, Y)]$ of any alternative forecast G .

A natural consistency criterion for S is that, for observations assumed to be drawn from F , it will not assess any other forecast G as being better than F itself, that is, that

$$\mathbb{E}_F[S(F, Y)] \leq \mathbb{E}_F[S(G, Y)] \quad (1)$$

for any F, G . A scoring rule meeting this criterion is called *proper*. If S were improper, then from the perspective of a forecaster focussed (solely) on expected loss minimization, the decision to state a forecast G other than the forecast F which they believe describes Y could be superior to the decision to state F . S is *strictly proper* when (1) is a strict inequality, in which case the *only* optimal decision for a forecaster seeking to minimize their expected loss is to state the forecast they believe to be true.

2.1 The allocation score is proper

Our primary decision theoretical procedure, outlined in section 2.2.1 of the main text, uses a decision problem with loss function $s(x, y)$ to define a scoring rule

$$S(F, y) := s(x^F, y) \quad (2)$$

where $x^F := \operatorname{argmin}_x \mathbb{E}_F[s(x, Y)]$ is the Bayes act for F with respect to s . Such scoring rules, which we call *Bayes scoring rules*, are proper by construction since

$$\begin{aligned} \mathbb{E}_F[S(F, Y)] &= \mathbb{E}_F[s(x^F, Y)] \\ &= \min_x \mathbb{E}_F[s(x, Y)] \quad (\text{by definition of } x^F) \end{aligned} \quad (3)$$

$$\begin{aligned} &\leq \mathbb{E}_F[s(x^G, Y)] \\ &= \mathbb{E}_F[S(G, Y)]. \end{aligned} \quad (4)$$

The allocation scoring rule is Bayes and therefore proper.

We note that in the probabilistic forecasting literature (see e.g., Gneiting [2011], Theorem 3) what we have termed Bayes scoring rules typically appear via (2) where x^F is some given functional of F which can be shown to be *elicitable*, that is, to be the Bayes act for some loss function s . Such a loss function is said to be a *consistent loss (or scoring) function* for the functional $F \mapsto x^F$, and many important recent results in the literature (e.g., Fissler and Ziegel [2016]) address whether there *exists* any loss function that is consistent for x^F . Our orientation is different from this insofar as we *begin* by specifying a decision problem and a loss function of subject matter relevance and use the Bayes act only as a bridge to a proper scoring rule. Consistency is never in doubt.

3 Expected shortages

A key feature of loss functions used in the decision theoretic definition of quantiles and related scoring rules such as the CRPS and the WIS (see ...), as well as the allocation loss function presented in this work, is the presence of a *shortage*: the amount $\max\{0, y - x\}$ by which a resource demand y exceeds a supply decision variable x , which, for convenience, we write as $(y - x)_+$. In particular, a quantile at probability level α of a distribution F on \mathbb{R}^1 (which we assume to have a well-defined density $f(x)$) is a Bayes act for the loss function

$$l(x, y) = Cx + L(y - x)_+$$

where $\alpha = 1 - C/L$ and C and L can be interpreted as the cost per unit of a resource (such as medicine) and the loss incurred when a unit of demand (such as illness) cannot be met due to the

shortage $(y - x)_+$. This follows because a Bayes act, as a minimizer of $\mathbb{E}_F[l(x, Y)]$, must also be a vanishing point of the derivative

$$\begin{aligned} \frac{d}{dx} \mathbb{E}_F[l(x, Y)] &= \mathbb{E}_F \left[\frac{d}{dx} l(x, Y) \right] \\ &= C + L \mathbb{E}_F \left[\frac{d}{dx} (Y - x)_+ \right] \\ &= C - L \mathbb{E}_F[\mathbf{1}\{Y > x\}] \\ &= C + L(F(x) - 1), \end{aligned} \quad (5)$$

so that $1 - C/L = F(x)$. The formula $\frac{d}{dx} \mathbb{E}_F[(Y - x)_+] = F(x) - 1$ for the derivative of the shortage, used above in (5), can be obtained from an application of the ‘‘Leibniz Rule’’:

$$\begin{aligned} \frac{d}{dx} \mathbb{E}_F[(Y - x)_+] &= \frac{d}{dx} \int_x^\infty (y - x) f_Y(y) dy \\ &= \int_x^\infty \frac{d}{dx} (y - x) f_Y(y) dy - (x - x) f_Y(x) = - \int_x^\infty f_Y(y) dy = F(x) - 1. \end{aligned} \quad (6)$$

Note that more care is required when F does not have a density. We will also use this result below in deriving the Bayes act for the allocation loss.

4 Allocation Bayes acts as vectors of marginal quantiles.

Here we study the form of the Bayes act for the allocation problem (AP) (equation (3) in section 2.2.2) of the text:

$$\underset{0 \leq x}{\text{minimize}} \mathbb{E}_F[s_A(x, Y)] = \sum_{i=1}^N L \cdot \mathbb{E}_{F_i}[(Y_i - x_i)_+] \text{ subject to } \sum_{i=1}^N x_i = K, \quad (7)$$

where the marginal forecasts F_i for $i = 1, \dots, N$ represent forecasts for N distinct locations. We show that the Bayes act $x^{F,K} = (x_1^{F,K}, \dots, x_N^{F,K})$ for a forecast F and resource constraint level K is a vector of quantiles of the marginal forecast distributions F_i at a single probability level $\tau^{F,K}$, that is, $x_i^{F,K} = q_{F_i, \tau^{F,K}}$. An immediate consequence used in the examples in Section 2.1 in the main text is that if $F_i = \text{Exp}(1/\sigma_i)$ for all i , then the Bayes act is proportional to $(\sigma_1, \dots, \sigma_N)$, since $q_{\text{Exp}(1/\sigma), \tau} = -\sigma \log(1 - \tau)$.

In order for $x^* \in \mathbb{R}_+^N$ to solve the AP it must be true that reallocating $\delta > 0$ units of the resource from location i to location j will lead to a net increase in expected shortage — in other words, the reallocation increases the expected shortage in location i is at least as much as it decreases the expected shortage in location j :

$$\begin{aligned} &\mathbb{E}_{F_i}[(Y_i - (x_i^* - \delta))_+] - \mathbb{E}_{F_i}[(Y_i - x_i^*)_+] \text{ (increase in } i) \\ &\geq \mathbb{E}_{F_j}[(Y_j - x_j^*)_+] - \mathbb{E}_{F_j}[(Y_j - (x_j^* + \delta))_+] \text{ (decrease in } j). \end{aligned}$$

Dividing by δ and letting $\delta \searrow 0$, this implies from (6) that

$$\begin{aligned} 1 - F_i(x_i^*) &= -\frac{d}{dx_i} \mathbb{E}_{F_i}[(Y_i - x_i^*)_+] \\ &= \lim_{\delta \searrow 0} \frac{1}{\delta} \{ \mathbb{E}_{F_i}[(Y_i - (x_i^* - \delta))_+] - \mathbb{E}_{F_i}[(Y_i - x_i^*)_+] \} \text{ (increase in } i) \\ &\geq \lim_{\delta \searrow 0} \frac{1}{\delta} \{ \mathbb{E}_{F_j}[(Y_j - x_j^*)_+] - \mathbb{E}_{F_j}[(Y_j - (x_j^* + \delta))_+] \} \text{ (decrease in } j) \\ &= -\frac{d}{dx_j} \mathbb{E}_{F_j}[(Y_j - x_j^*)_+] = 1 - F_j(x_j^*) \end{aligned} \quad (8)$$

Note that negative derivatives appear because our optimality condition addresses how a *decrease* in resources will *increase* the expected shortage in i and vice versa in j . Since (8) also holds with i and j

reversed, a number λ (a *Lagrange multiplier*) exists such that $L(1 - F_k(x_k^*)) = \lambda$ for all $k \in 1, \dots, N$. (We scale by L to facilitate possible future interpretations of λ in terms of the partial derivatives of $\mathbb{E}_F[s_A(x, Y)]$.) That is, x_k^* is a quantile q_{τ, F_k} for $\tau = 1 - \lambda/L$. The value of τ is then determined by the constraint equation

$$\sum_{i=1}^N q_{\tau, F_i} = K. \quad (9)$$

It is important to note that τ depends on F and K and is *not* a fixed parameter of the allocation scoring rule.

5 Numerical computation of allocation Bayes acts

To compute an allocation score $S_A(F, y; K) := s_A(x^{F, K}, y)$, we require numerical values for a Bayes act solving the AP (7) — that is, we must find the specific resource allocations for each location that are determined by the forecast F under the resource constraint K . Assuming we have reliable means of calculating quantiles q_{α, F_i} of the marginal forecasts F_i , these allocations are given by q_{τ^*, F_i} where τ^* solves the equation (9). However, this equation is not analytically tractable and we must resort to a numerical method for finding an approximation $\tilde{\tau}$ of τ^* .

We have implemented an iterative bisection method that makes use of the fact that $\sum_{i=1}^N q_{\tau, F_i}$ is an increasing function of τ . The algorithm begins with an initial search interval $[\tau_{L,1}, \tau_{U,1}]$ (such as $[0, \max_i F_i(K)]$) that clearly contains the solution τ^* . At each step j of the algorithm, we evaluate the total allocation $\sum_{i=1}^N q_{\tau_{M,j}, F_i}$ at the midpoint of the search interval, $\tau_{M,j} = \frac{1}{2}(\tau_{L,j} + \tau_{U,j})$ and continue the search on the narrowed sub-interval

$$[\tau_{L,j+1}, \tau_{U,j+1}] = \begin{cases} [\tau_{L,j}, \tau_{M,j}] & \text{if } \sum_{i=1}^N q_{\tau_{M,j}, F_i} \geq K \\ [\tau_{M,j}, \tau_{U,j}] & \text{if } \sum_{i=1}^N q_{\tau_{M,j}, F_i} < K. \end{cases}$$

This search continues until $\tau_{U,j+1} < (1 + \varepsilon)\tau_{L,j+1}$ for a suitably small $\varepsilon > 0$. We have implemented this procedure along with the resulting score computations in the R package `alloscore` which provided all allocation score values used in the analysis of section 3 in the main text.

Subtleties can arise when the forecast densities f_i vanish or are very small, in which case quantiles are non-unique or highly variable near a probability level, leading to ambiguity or numerical instabilities in the evaluation of $\sum_{i=1}^N q_{\tau, F_i}$. Additionally, if point masses are present in any of the F_i , (9) will not have a unique solution for some discrete set of constraint levels K . We have adopted conventions for detecting such levels and enforcing consistency in score calculations near them. Through extensive experimentation, we have determined that these conditions seem to address these challenges with the forecasts we are working with, but we leave a more rigorous approximation error analysis for later work.

6 Computing allocations from finite quantile forecast representations

In section 3 of the manuscript, we used the allocation score to evaluate forecasts of COVID-19 hospitalizations that have been submitted to the US COVID-19 Forecast Hub. These forecasts are submitted to the Hub using a set of 23 quantiles of the forecast distribution at the 23 probability levels in the set $\mathcal{T} = \{0.01, 0.025, 0.05, 0.1, 0.15, \dots, 0.9, 0.95, 0.975, 0.99\}$, which specify a predictive median and the endpoints of central $(1 - \alpha) \times 100\%$ prediction intervals at levels $\alpha = 0.02, 0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9$. For a given week and target date, we use $q_{i,k}$ to denote the submitted quantiles for location i and probability level $\tau_k \in \mathcal{T}$, $k = 1, \dots, 23$.

In the event that there is some $k \in \{1, \dots, 23\}$ for which $\sum_i q_{i,k} = K$, i.e., the provided predictive quantiles at level τ_k sum across locations to the resource constraint K , the solution to the allocation problem is given by those quantiles. However, generally this will not be the case; the optimal allocation will typically be at some probability level $\tau^* \notin \mathcal{T}$.

To address this situation and support the numerical allocation algorithm outlined in section 5, we need a mechanism to approximate the full cumulative distribution functions F_i , $i = 1, \dots, N$ based on the provided quantiles. We have developed functionality for this purpose in the `distfromq` package for R.[Ray et al., 2023] This functionality represents a distribution as a mixture of discrete and continuous parts, and it works in two steps:

1. Identify a discrete component of the distribution consisting of zero or more point masses, and create an adjusted set of predictive quantiles for the continuous part of the distribution by subtracting the point mass probabilities and rescaling.
2. For the continuous part of the distribution, different approaches are used on the interior and exterior of the provided quantiles:
 - (a) On the interior, a monotonic cubic spline interpolates the adjusted quantiles representing the continuous part of the distribution.
 - (b) A location-scale parametric family is used to extrapolate beyond the provided quantiles. The location and scale parameters are estimated separately for the lower and upper tails so as to obtain a tail distribution that matches the two most extreme quantiles in each tail.

The resulting distributional estimate exactly matches all of the predictive quantiles provided by the forecaster. We use the cumulative distribution function resulting from this procedure as an input to the allocation score algorithm.

We refer the reader to the `distfromq` documentation for further detail.[Ray et al., 2023]

7 Propriety of parametric approximation

In section 2.1 we saw that Bayes scoring rules (such as the allocation score) are proper. However, this result implicitly assumes a perfect description of the full forecast distribution F is available for computing the Bayes act. In practice, open forecasting exercises are not able to collect a description of the forecast distribution F that suffices for the computation of Bayes acts for arbitrary loss functions other than in simple settings such as the elicitation of probabilities for all levels of a categorical variable. In settings where the outcome being forecasted is a continuous quantity (such as the proportion of outpatient doctor visits where the patient had influenza-like illness) or a large count (such as influenza hospitalizations), forecasting exercises instead collect low-dimensional forecast distribution summaries, such as the 23 quantiles (see section 6) collected by the US COVID-19 Forecast Hub.

Bayes scoring rules are invariant under reparametrization of the action space.

Let $S(F, y)$ be a Bayes scoring rule derived from an action set \mathcal{X} and loss $s(x, y)$. Let \mathcal{W} be another set and $M : \mathcal{W} \rightarrow \mathcal{X}$ a map. M defines a loss $s_M(\theta, y) := s(M(\theta), y)$ on \mathcal{W} . This in turn defines the Bayes scoring rule

$$S_M(F, y) := s_M(\theta^F, y) = s(M(\theta^F), y). \quad (10)$$

If M is surjective onto the set of Bayes acts $\{x \in \mathcal{X} \mid x = x^F \text{ for some } F \in \mathcal{P}\}$, then the scores coincide since

$$S_M(F, y) = s_M(\theta^F, y) = \min_{\theta} \{s_M(\theta, F)\} \quad (11)$$

$$= \min_{\theta} \{s(M(\theta), F)\} \quad (12)$$

$$= \min_x \{s(x, F)\} = s(x^F, F) = S(F, y). \quad (13)$$

(Note that this does not imply that $M(\theta^F) = x^F$ since x^F might not be a unique Bayes act.)

Identifying $G_{\theta} \in \mathcal{P}_{\text{dfq}}$ with its quantile parameters θ , the map $M(\theta) := x^{\theta}$ is surjective onto the action set of allocations under any constraint K since they can be specified as quantiles for distributions in \mathcal{P}_{dfq} . Therefore the `distfromq` “reparametrized” allocation scoring rule $S_{A, \text{dfq}}(F, y) := s_{\text{dfq}}(\theta^F, y) = s_A(x^{\theta^F}, y)$ coincides with the AS $S_A(F, y)$ itself.

Unfortunately,

For convenience we write $S(F, G) := \mathbb{E}_G[S(F, Y)]$. Let

$$\begin{aligned} q_{23}(F) &:= (q_{.01, F}, \dots, q_{.99, F}) \\ G^q(F) &:= G_{q_{23}(F)} \in \mathcal{P}_{\text{dfq}} \\ G^*(F) &:= \operatorname{argmin}_{G \in \mathcal{P}_{\text{dfq}}} S_A(G, F) \end{aligned}$$

We have the “right” approximate allocation score

$$S_R(F, y) := S_A(G^*(F), y)$$

which is proper because

$$\begin{aligned} S_R(F, F) &= S_A(G^*(F), F) \\ &\leq S_A(G^*(F), H) \quad (\text{by definition of } G^*) \\ &= S_R(F, H). \end{aligned}$$

This seems to almost be what is called “properization” in Brehmer and Gneiting [2020]. We also have a “wrong” approximation

$$\tilde{S}_R(F, y) := S_A(G^q(F), y)$$

which is improper because

$$\begin{aligned} \tilde{S}_R(F, G^q(F)) &= S_A(G^q(F), G^q(F)) \\ &< S_A(G^q(F), F) \quad (\text{because } S_A \text{ is proper}) \\ &= \tilde{S}_R(F, F). \end{aligned}$$

References

- Tilmann Gneiting. Making and evaluating point forecasts. *Journal of the American Statistical Association*, 106(494):746–762, 2011.
- Tobias Fissler and Johanna F. Ziegel. Higher order elicibility and Osband’s principle. *The Annals of Statistics*, 44(4):1680 – 1707, 2016.
- Evan L Ray, Aaron Gerding, and Li Shandross. *distfromq: Reconstruct a Distribution from a Collection of Quantiles*, 2023. URL <https://github.com/reichlab/distfromq>. R package version 1.0.2.
- Jonas R Brehmer and Tilmann Gneiting. Properization: constructing proper scoring rules via bayes acts. *Annals of the Institute of Statistical Mathematics*, 72:659–673, 2020.